

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

train


In [2]:

```
train_df=pd.read_csv(r"C:\Users\shaha\OneDrive\Desktop\Excel\Mobile_Price_Classification.
train_df
```

Out[2]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_
0	842	0	2.2	0	1	0	7	0.6	18
1	1021	1	0.5	1	0	1	53	0.7	18
2	563	1	0.5	1	2	1	41	0.9	14
3	615	1	2.5	0	0	0	10	0.8	18
4	1821	1	1.2	0	13	1	44	0.6	14
...
1995	794	1	0.5	1	0	1	2	0.8	10
1996	1965	1	2.6	1	0	0	39	0.2	18
1997	1911	0	0.9	1	1	1	36	0.7	10
1998	1512	0	0.9	0	4	1	46	0.1	14
1999	510	1	2.0	1	5	1	45	0.9	16

2000 rows × 21 columns



In [3]:

```
train_df.head()
```

Out[3]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt
0	842	0	2.2	0	1	0	7	0.6	188
1	1021	1	0.5	1	0	1	53	0.7	136
2	563	1	0.5	1	2	1	41	0.9	145
3	615	1	2.5	0	0	0	10	0.8	131
4	1821	1	1.2	0	13	1	44	0.6	141

5 rows × 21 columns




In [4]:

```
train_df.tail()
```

Out[4]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_v
1995	794	1	0.5	1	0	1	2	0.8	10
1996	1965	1	2.6	1	0	0	39	0.2	18
1997	1911	0	0.9	1	1	1	36	0.7	10
1998	1512	0	0.9	0	4	1	46	0.1	14
1999	510	1	2.0	1	5	1	45	0.9	16

5 rows × 21 columns



In [5]:

```
train_df.shape
```

Out[5]:

```
(2000, 21)
```

In [6]:

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   battery_power       2000 non-null   int64
1   blue                 2000 non-null   int64
2   clock_speed         2000 non-null   float64
3   dual_sim            2000 non-null   int64
4   fc                   2000 non-null   int64
5   four_g              2000 non-null   int64
6   int_memory          2000 non-null   int64
7   m_dep               2000 non-null   float64
8   mobile_wt           2000 non-null   int64
9   n_cores             2000 non-null   int64
10  pc                   2000 non-null   int64
11  px_height            2000 non-null   int64
12  px_width             2000 non-null   int64
13  ram                  2000 non-null   int64
14  sc_h                 2000 non-null   int64
15  sc_w                 2000 non-null   int64
16  talk_time            2000 non-null   int64
17  three_g              2000 non-null   int64
18  touch_screen         2000 non-null   int64
19  wifi                 2000 non-null   int64
20  price_range          2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
```

In [7]:

```
train_df.describe()
```

Out[7]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_m
count	2000.000000	2000.0000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	1238.518500	0.4950	1.522250	0.509500	4.309500	0.521500	32.000000
std	439.418206	0.5001	0.816004	0.500035	4.341444	0.499662	18.000000
min	501.000000	0.0000	0.500000	0.000000	0.000000	0.000000	2.000000
25%	851.750000	0.0000	0.700000	0.000000	1.000000	0.000000	16.000000
50%	1226.000000	0.0000	1.500000	1.000000	3.000000	1.000000	32.000000
75%	1615.250000	1.0000	2.200000	1.000000	7.000000	1.000000	48.000000
max	1998.000000	1.0000	3.000000	1.000000	19.000000	1.000000	64.000000

8 rows × 21 columns



In [8]:

```
x=train_df.drop('wifi',axis=1)
y=train_df['wifi']
```

In [9]:

```
train_df['clock_speed'].value_counts()
```

```
0.6    74
1.4    70
1.3    68
1.5    67
2.0    67
1.9    65
0.7    64
2.9    62
1.8    62
1.0    61
1.7    60
2.2    59
0.9    58
2.4    58
0.8    58
1.2    56
2.6    55
2.7    55
1.1    51
3.0    28
```

In [10]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7,random_state=25)
x_train.shape
```

Out[10]:

(1400, 20)

In [11]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[11]:

```
RandomForestClassifier()
```

In [12]:

```
rf=RandomForestClassifier()
```


In [13]:

```
params={'max_depth':[2,3,5,10,20],  
        'min_samples_leaf':[5,10,20,50,100,200],  
        'n_estimators':[10,25,30,50,100,200]}
```

In [14]:

```
from sklearn.model_selection import GridSearchCV  
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')  
grid_search.fit(x_train,y_train)
```

Out[14]:



In [15]:

```
grid_search.best_score_
```

Out[15]:

0.5178571428571428

In [16]:

```
rf_best=grid_search.best_estimator_  
print(rf_best)
```


RandomForestClassifier(max_depth=2, min_samples_leaf=50, n_estimators=25)

In [17]:

```
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=['1','0'],filled=True)
```

Out[17]:

```
[Text(0.4, 0.8333333333333334, 'px_height <= 109.5\ngini = 0.5\nsamples = 889\nvalue = [718, 682]\nclass = 1'),
Text(0.2, 0.5, 'gini = 0.429\nsamples = 81\nvalue = [93, 42]\nclass = 1'),
Text(0.6, 0.5, 'px_width <= 670.5\ngini = 0.5\nsamples = 808\nvalue = [625, 640]\nclass = 0'),
Text(0.4, 0.16666666666666666, 'gini = 0.458\nsamples = 75\nvalue = [45, 82]\nclass = 0'),
Text(0.8, 0.16666666666666666, 'gini = 0.5\nsamples = 733\nvalue = [580, 558]\nclass = 1')]
```



In [18]:

```
rf_best.feature_importances_
```

Out[18]:

```
array([0.0346861 , 0.01871166, 0.01708762, 0.          , 0.12656838,
        0.          , 0.11360976, 0.04030123, 0.05479728, 0.02367189,
        0.03870892, 0.14946964, 0.1034957 , 0.06271748, 0.03405665,
        0.01071626, 0.12214479, 0.          , 0.02576145, 0.0234952 ])
```

In [19]:

```
imp_df=pd.DataFrame({'Varname':x_train.columns,"Imp":rf_best.feature_importances_})  
imp_df.sort_values(by="Imp",ascending=False)
```

Out[19]:

	Varname	Imp
11	px_height	0.149470
4	fc	0.126568
16	talk_time	0.122145
6	int_memory	0.113610
12	px_width	0.103496
13	ram	0.062717
8	mobile_wt	0.054797
7	m_dep	0.040301
10	pc	0.038709
0	battery_power	0.034686
14	sc_h	0.034057
18	touch_screen	0.025761
9	n_cores	0.023672
19	price_range	0.023495
1	blue	0.018712
2	clock_speed	0.017088
15	sc_w	0.010716
5	four_g	0.000000
3	dual_sim	0.000000
17	three_g	0.000000

test

In [20]:

```
test_df=pd.read_csv(r"C:\Users\shaha\OneDrive\Desktop\Excel\Mobile_Price_Classification_
test_df
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	p
0	1	1043	1	1.8	1	14	0	5	0.1	193	...	1
1	2	841	1	0.5	1	4	1	61	0.8	191	...	1
2	3	1807	1	2.8	0	1	0	27	0.9	186	...	
3	4	1546	0	0.5	1	18	1	25	0.5	96	...	2
4	5	1434	0	1.4	0	11	1	49	0.5	108	...	1
...	
995	996	1700	1	1.9	0	0	1	54	0.5	170	...	1
996	997	609	0	1.8	1	0	0	13	0.9	186	...	
997	998	1185	0	1.4	0	1	1	8	0.5	80	...	1
998	999	1533	1	0.5	1	0	0	50	0.4	171	...	1
999	1000	1270	1	0.5	0	4	1	35	0.1	140	...	1

In [21]:

```
test_df.head()
```

Out[21]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_
0	1	1043	1	1.8	1	14	0	5	0.1	1
1	2	841	1	0.5	1	4	1	61	0.8	1
2	3	1807	1	2.8	0	1	0	27	0.9	1
3	4	1546	0	0.5	1	18	1	25	0.5	
4	5	1434	0	1.4	0	11	1	49	0.5	1

5 rows × 21 columns

In [22]:

test_df.tail()

Out[22]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mo
995	996	1700	1	1.9	0	0	1	54	0.5	
996	997	609	0	1.8	1	0	0	13	0.9	
997	998	1185	0	1.4	0	1	1	8	0.5	
998	999	1533	1	0.5	1	0	0	50	0.4	
999	1000	1270	1	0.5	0	4	1	35	0.1	

5 rows × 21 columns

In [23]:

test_df.shape

Out[23]:

(1000, 21)

In [24]:

test_df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   id                     1000 non-null   int64  
1   battery_power          1000 non-null   int64  
2   blue                   1000 non-null   int64  
3   clock_speed            1000 non-null   float64 
4   dual_sim               1000 non-null   int64  
5   fc                     1000 non-null   int64  
6   four_g                 1000 non-null   int64  
7   int_memory             1000 non-null   int64  
8   m_dep                  1000 non-null   float64 
9   mobile_wt              1000 non-null   int64  
10  n_cores                 1000 non-null   int64  
11  pc                      1000 non-null   int64  
12  px_height               1000 non-null   int64  
13  px_width               1000 non-null   int64  
14  ram                     1000 non-null   int64  
15  sc_h                    1000 non-null   int64  
16  sc_w                    1000 non-null   int64  
17  talk_time               1000 non-null   int64  
18  three_g                 1000 non-null   int64  
19  touch_screen            1000 non-null   int64  
20  wifi                    1000 non-null   int64  
dtypes: float64(2), int64(19)
memory usage: 164.2 KB

```


In [25]:

```
test_df.describe()
```

Out[25]:

	id	battery_power	blue	clock_speed	dual_sim	fc
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	500.500000	1248.510000	0.516000	1.540900	0.517000	4.593000
std	288.819436	432.458227	0.499994	0.829268	0.499961	4.463325
min	1.000000	500.000000	0.000000	0.500000	0.000000	0.000000
25%	250.750000	895.000000	0.000000	0.700000	0.000000	1.000000
50%	500.500000	1246.500000	1.000000	1.500000	1.000000	3.000000
75%	750.250000	1629.250000	1.000000	2.300000	1.000000	7.000000
max	1000.000000	1999.000000	1.000000	3.000000	1.000000	19.000000

8 rows × 21 columns



In [26]:

```
x=test_df.drop('blue',axis=1)
y=test_df['blue']
```

In [27]:

```
test_df['blue'].value_counts()
```

Out[27]:

```
blue
1    516
0    484
Name: count, dtype: int64
```

In [28]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7)
x_test.shape
```

Out[28]:

(300, 20)

In [29]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_test,y_test)
```

Out[29]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In [30]:

```
rf=RandomForestClassifier()
```

In [31]:

```
params={'max_depth':[2,3,5,10,20],
        'min_samples_leaf':[5,10,20,50,100,200],
        'n_estimators':[10,25,30,50,100,200]}
```

In [32]:

```
from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')
grid_search.fit(x_test,y_test)
```

Out[32]:

```
► GridSearchCV
► estimator: RandomForestClassifier
  ► RandomForestClassifier
```

In [33]:

```
grid_search.best_score_
```

Out[33]:

```
0.5866666666666667
```

In [34]:

```
rf_best=grid_search.best_estimator_
print(rf_best)
```

```
RandomForestClassifier(max_depth=10, min_samples_leaf=5, n_estimators=50)
```

In [35]:

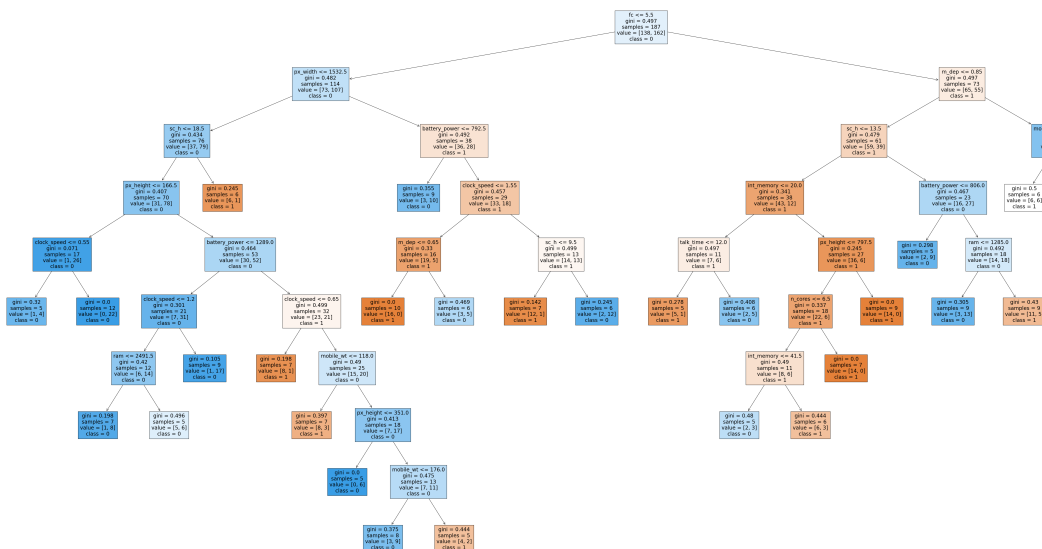
```
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=['1','0'],filled=True)
```

Out[35]:

```

[Text(0.5703125, 0.95, 'fc <= 5.5\ngini = 0.497\nsamples = 187\nvalue = [1
38, 162]\nclasse = 0'),
Text(0.2890625, 0.85, 'px_width <= 1532.5\ngini = 0.482\nsamples = 114\nv
alue = [73, 107]\nclasse = 0'),
Text(0.171875, 0.75, 'sc_h <= 18.5\ngini = 0.434\nsamples = 76\nvalue =
[37, 79]\nclasse = 0'),
Text(0.140625, 0.65, 'px_height <= 166.5\ngini = 0.407\nsamples = 70\nval
ue = [31, 78]\nclasse = 0'),
Text(0.0625, 0.55, 'clock_speed <= 0.55\ngini = 0.071\nsamples = 17\nvalu
e = [1, 26]\nclasse = 0'),
Text(0.03125, 0.45, 'gini = 0.32\nsamples = 5\nvalue = [1, 4]\nclasse =
0'),
Text(0.09375, 0.45, 'gini = 0.0\nsamples = 12\nvalue = [0, 22]\nclasse =
0'),
Text(0.21875, 0.55, 'battery_power <= 1289.0\ngini = 0.464\nsamples = 53
\nvalue = [30, 52]\nclasse = 0'),
Text(0.15625, 0.45, 'clock_speed <= 1.2\ngini = 0.301\nsamples = 21\nvalu
e = [7, 31]\nclasse = 0'),
Text(0.125, 0.35, 'ram <= 2491.5\ngini = 0.42\nsamples = 12\nvalue = [6,
14]\nclasse = 0'),
Text(0.09375, 0.25, 'gini = 0.198\nsamples = 7\nvalue = [1, 8]\nclasse =
0'),
Text(0.15625, 0.25, 'gini = 0.496\nsamples = 5\nvalue = [5, 6]\nclasse =
0'),
Text(0.1875, 0.35, 'gini = 0.105\nsamples = 9\nvalue = [1, 17]\nclasse =
0'),
Text(0.28125, 0.45, 'clock_speed <= 0.65\ngini = 0.499\nsamples = 32\nval
ue = [23, 21]\nclasse = 1'),
Text(0.25, 0.35, 'gini = 0.198\nsamples = 7\nvalue = [8, 1]\nclasse = 1'),
Text(0.3125, 0.35, 'mobile_wt <= 118.0\ngini = 0.49\nsamples = 25\nvalue
= [15, 20]\nclasse = 0'),
Text(0.28125, 0.25, 'gini = 0.397\nsamples = 7\nvalue = [8, 3]\nclasse =
1'),
Text(0.34375, 0.25, 'px_height <= 351.0\ngini = 0.413\nsamples = 18\nvalu
e = [7, 17]\nclasse = 0'),
Text(0.3125, 0.15, 'gini = 0.0\nsamples = 5\nvalue = [0, 6]\nclasse = 0'),
Text(0.375, 0.15, 'mobile_wt <= 176.0\ngini = 0.475\nsamples = 13\nvalue
= [7, 11]\nclasse = 0'),
Text(0.34375, 0.05, 'gini = 0.375\nsamples = 8\nvalue = [3, 9]\nclasse =
0'),
Text(0.40625, 0.05, 'gini = 0.444\nsamples = 5\nvalue = [4, 2]\nclasse =
1'),
Text(0.203125, 0.65, 'gini = 0.245\nsamples = 6\nvalue = [6, 1]\nclasse =
1').

```



```

0'),
In[36]: rf_best.feature_importances_
Out[36]: array([0.0553984, 0.07385643, 0.04627033, 0.01135697, 0.05588569,
0.625, 0.55, 0.01024495, 0.05718877, 0.06860488, 0.06803284, 0.02793173,
[7, 6]\nnclass = 1'),
Text(0.58375, 0.45, 'gini = 0.11880176, 0.06892912, 0.06977885, 0.06849571,
1'),
Text(0.65625, 0.45, 'gini = 0.408\nnsamples = 6\nnvalue = [2, 5]\nnclass =
0'),
In[37]:
Text(0.75, 0.55, 'px height <= 797.5\nngini = 0.245\nnsamples = 27\nnvalue =
130, 6]\nnclass = 1'),
In[38]: imp_df = pd.DataFrame({'Varname':x_train.columns,"Imp":rf_best.feature_importances_})
imp_df.sort_values(by="Imp",ascending=False)
Out[38]:
Text(0.71875, 0.45, 'n_cores <= 0.5\nngini = 0.337\nnsamples = 18\nnvalue =
[22, 6]\nnclass = 1'),
Out[37]:
Text(0.6875, 0.35, 'int_memory <= 41.5\nngini = 0.49\nnsamples = 11\nnvalue
= [8, 6]\nnclass = 1'),
Text(0.65625, 0.25, 'Imp', 'gini = 0.48\nnsamples = 5\nnvalue = [2, 3]\nnclass =
0'),
Text(0.71875, 0.25, 'px height <= 1180.2\nngini = 0.444\nnsamples = 6\nnvalue = [6, 3]\nnclass =
1'),
Text(0.75, 0.35, 'gini = 0.0\nnsamples = 7\nnvalue = [14, 0]\nnclass = 1'),
Text(0.78125, 0.45, 'gini = 0.0\nnsamples = 9\nnvalue = [14, 0]\nnclass =
1'),
Text(0.84375, 0.65, 'battery_power <= 806.0\nngini = 0.467\nnsamples = 23\nn
value = [16, 27]\nnclass = 0'),
Text(0.8125, 0.55, 'gini = 0.298\nnsamples = 5\nnvalue = [2, 9]\nnclass =
0'),
Text(0.875, 0.55, 'ram <= 1285.0\nngini = 0.492\nnsamples = 18\nnvalue = [1
4, 18]\nnclass = 0'),
Text(0.84375, 0.45, 'gini = 0.305\nnsamples = 9\nnvalue = [3, 13]\nnclass =
0'),
Text(0.90625, 0.45, 'gini = 0.43\nnsamples = 9\nnvalue = [11, 5]\nnclass =
1'),
Text(0.9375, 0.75, 'mobile_wt <= 111.5\nngini = 0.397\nnsamples = 12\nnvalue
= [6, 16]\nnclass = 0'),
Text(0.90625, 0.65, 'gini = 0.5\nnsamples = 6\nnvalue = [6, 6]\nnclass =
1'),
Text(0.96875, 0.65, 'gini = 0.0\nnsamples = 6\nnvalue = [0, 10]\nnclass =
0'),
9      n_cores  0.027932
18  touch_screen  0.021275
19      wifi  0.013383
3      dual_sim  0.011357
17      three_g  0.010981
5      four_g  0.010245

```